



SCIT

**School of Computing and
Information Technology**

CSCI317

Database Performance Tuning

This paper is for students studying at the Singapore Institute of Management Pte Ltd.

S3-2020 FINAL EXAMINATION

Date: ???

Time: ???

Exam value: **40% of the subject assessment**

Marks available: **40 marks**

DIRECTIONS TO CANDIDATES

- (1) The answers to the questions included in the final examination must be hand written with a BLACK or DARK BLUE PEN on the WHITE PIECES of paper in A4 format. No pencil and no other colour of paper is allowed.
- (2) When finished, take the pictures of the hand-written solution, save the pictures in files (jpeg, jpg, gif, bmp, png formats are all acceptable), and submit the files through Moodle. Using mobile phone cameras is all right. It is possible to take more than one picture per answer to assure the good readability of an answer. The marks will be deducted for submissions in the different formats. No more than 20 files can be submitted and no more than 200Mbytes can be submitted. Please well plan your pictures.
- (3) The file must have the names indicating a number of the respective question in the final examination paper like q1, q2, ... and q1-1, q1-2, ... when more than one picture is used for an answer of a question. Marks will be deducted for the incorrect file names.
- (4) All answers including the drawings must be hand written. No printed material will be evaluated.
- (5) Marks will be deducted for the late submissions at a rate of 1 mark per 1 minute late.

Introduction

The questions 2, 4, 5, and 6 of the examination paper are related to the following simplified version of TPC-H benchmark database used in the laboratory classes.

```
CUSTOMER (
C_CUSTKEY      NUMBER(12)      NOT NULL,
C_NAME         VARCHAR(25)     NOT NULL,
C_ADDRESS      VARCHAR(40)     NOT NULL,
C_NATIONKEY    NUMBER(12)     NOT NULL,
C_ACCTBAL      NUMBER(6)       NOT NULL,
C_PHONE        NUMBER(12)     NOT NULL,
CONSTRAINT CUSTOMER_PKEY PRIMARY KEY(C_CUSTKEY) );

PART (
P_PARTKEY      NUMBER(12)      NOT NULL,
P_NAME         VARCHAR(55)     NOT NULL,
P_BRAND        CHAR(10)        NOT NULL,
P_SIZE         NUMBER(12)     NOT NULL,
P_RETAILPRICE  NUMBER(12,2)   NOT NULL,
CONSTRAINT PART_PKEY PRIMARY KEY (P_PARTKEY) );

PARTSUPP (
PS_PARTKEY     NUMBER(12)      NOT NULL,
PS_SUPPNAME    VARCHAR(55)     NOT NULL,
PS_AVAILQTY    NUMBER(12)     NOT NULL,
CONSTRAINT PARTSUPP_PKEY PRIMARY KEY (PS_PARTKEY, PS_SUPPNAME),
CONSTRAINT PARTSUPP_FKEY FOREIGN KEY (PS_PARTKEY)
REFERENCES PART(P_PARTKEY) );

ORDERS (
O_ORDERKEY     NUMBER(12)      NOT NULL,
O_CUSTKEY      NUMBER(12)     NOT NULL,
O_TOTALPRICE   NUMBER(12,2)   NOT NULL,
O_ORDERDATE    DATE           NOT NULL,
CONSTRAINT ORDERS_PKEY PRIMARY KEY (O_ORDERKEY),
CONSTRAINT ORDERS_FKEY1 FOREIGN KEY (O_CUSTKEY)
REFERENCES CUSTOMER(C_CUSTKEY) );

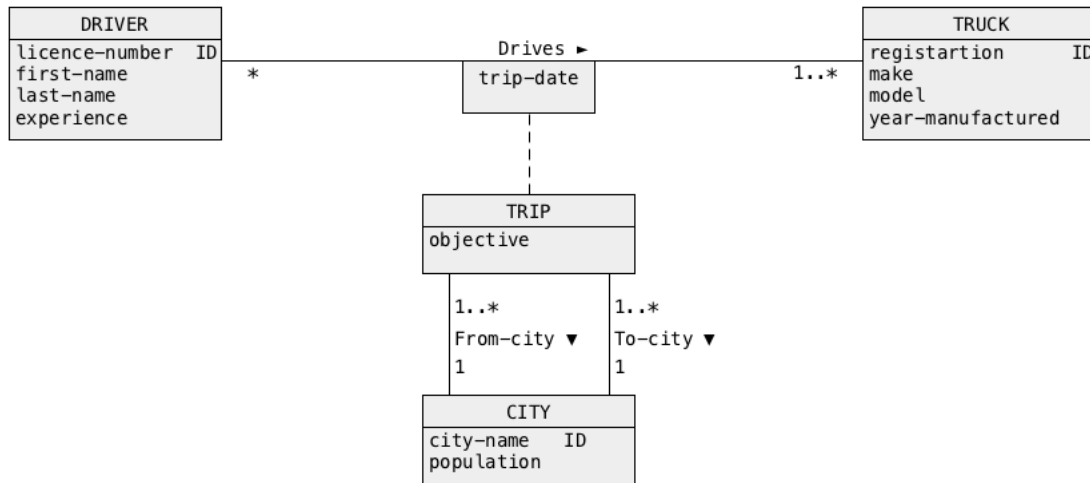
LINEITEM (
L_ORDERKEY     NUMBER(12)      NOT NULL,
L_PARTKEY      NUMBER(12)     NOT NULL,
L_LINENUMBER   NUMBER(12)     NOT NULL,
L_QUANTITY     NUMBER(12,2)   NOT NULL,
L_SHIPDATE     DATE           NOT NULL,
L_TAX          NUMBER(4,2)     NOT NULL,
CONSTRAINT LINEITEM_PKEY PRIMARY KEY (L_ORDERKEY, L_LINENUMBER),
CONSTRAINT LINEITEM_FKEY1 FOREIGN KEY (L_ORDERKEY)
REFERENCES ORDERS(O_ORDERKEY),
CONSTRAINT LINEITEM_FKEY2 FOREIGN KEY (L_PARTKEY)
REFERENCES PART(P_PARTKEY) );
```

Assume that, the relational tables listed above occupy the following amounts of disk storage:

CUSTOMER	100 Mbytes
PART	40 Mbytes
PARTSUPP	100 Mbytes
ORDERS	200 Mbytes
LINEITEM	700 Mbytes

Question 1**(7 marks)**

The following conceptual schema represents a database domain where the drivers use the trucks for the trips from city to city. We assume that a driver can make at most one trip per day. Each trip has an objective, like for example delivery of the ordered items, collection of parcels to be delivered to another place, etc. All other attributes are self-explanatory.



- (1) Perform simplification of the conceptual schema above and re-draw the simplified conceptual schema.

(2 marks)

- (2) We would like to improve the performance of the following class of applications:

Find the first and the last names of drivers (attributes `first-name`, `last-name` in a class `DRIVER`) who travelled between two given cities (attribute `city-name` in a class `CITY`) and used a vehicle manufactured before a given date (attribute `year-manufactured` in a class `TRUCK`).

The following application belongs to the class of applications given above.

Find the first and the last names of drivers who travelled from `Dapto` to `Sydney` and used an old vehicle manufactured before a year 2000.

Find the denormalizations of the simplified conceptual schema that improves the performance of the class of applications described above. Re-draw the simplified conceptual schema after the denormalizations.

(5 marks)

Question 2**(7 marks)**

Consider the following fragment of query processing plan.

Id	Operation	Name	Rows	Bytes	TempSpc	Cost (%CPU)	Time
0	SELECT STATEMENT		317K	39M		17968 (1)	00:00:01
* 1	HASH JOIN		317K	39M		17968 (1)	00:00:01
* 2	TABLE ACCESS FULL	CUSTOMER	40091	430K		390 (1)	00:00:01
* 3	HASH JOIN RIGHT ANTI		317K	35M	3080K	17577 (1)	00:00:01
* 4	TABLE ACCESS FULL	LINEITEM	150K	1318K		12153 (1)	00:00:01
* 5	TABLE ACCESS FULL	ORDERS	450K	46M		2698 (1)	00:00:01

Predicate Information (identified by operation id):

- 1 - access("O_CUSTKEY"="C_CUSTKEY")
- 2 - filter("C_ACCTBAL">200)
- 3 - access("O_ORDERKEY"="L_ORDERKEY")
- 4 - filter("L_TAX">0.1)
- 5 - filter("O_CUSTKEY">=0)

- (1) Find and draw a syntax tree of the query processing plan listed above. To draw a syntax tree, use the relational algebra operations (and NOT Oracle query processing plan operations) explained during the lecture classes.

(3 marks)

- (2) Discover and write SELECT statement that may have a query processing plan listed above.

(4 marks)

```
SQL> -- Explaining plan for the given query processing plan
SQL> explain plan for
2 select *
3 from orders join customer
4 on o_custkey = c_custkey
5 where o_orderkey not in
6 (select l_orderkey from lineitem where l_tax > 0.1)
7 and o_custkey >= 0
8 and c_acctbal > 200;
```

Explained.

```
SQL>
SQL> -- Displaying the explained plan
SQL> SELECT * FROM TABLE(DBMS_XPLAN.DISPLAY);
```

PLAN_TABLE_OUTPUT

Plan hash value: 1426367367

Id	Operation	Name	Rows	Bytes	TempSpc	Cost (%CPU)	Time
0	SELECT STATEMENT		379K	100M		20967 (1)	00:00:01
* 1	HASH JOIN		379K	100M	6696K	20967 (1)	00:00:01
* 2	TABLE ACCESS FULL	CUSTOMER	40091	6225K		390 (1)	00:00:01
* 3	HASH JOIN RIGHT ANTI		379K	42M	1536K	17917 (1)	00:00:01
* 4	TABLE ACCESS FULL	LINEITEM	74757	657K		12153 (1)	00:00:01
* 5	TABLE ACCESS FULL	ORDERS	450K	46M		3113 (1)	00:00:01

Predicate Information (identified by operation id):

- 1 - access("O_CUSTKEY"="C_CUSTKEY")
- 2 - filter("CUSTOMER"."C_ACCTBAL">200 AND "C_CUSTKEY">=0)
- 3 - access("ORDERS"."O_ORDERKEY"="L_ORDERKEY")
- 4 - filter("L_TAX">0.1)
- 5 - filter("ORDERS"."O_CUSTKEY">=0)

21 rows selected.

```
-- Explaining plan for the given query processing plan
explain plan for
select *
from orders join customer
on o_custkey = c_custkey
where o_orderkey not in
(select l_orderkey from lineitem where l_tax > 0.1)
and o_custkey >= 0
and c_acctbal > 200;

-- Displaying the explained plan
SELECT * FROM TABLE(DBMS_XPLAN.DISPLAY);

SPOOL OFF
```

Question 3**(6 marks)**

A relational table PARTSUPP contains information about the part supplied by suppliers.

PARTSUPP(supplier#, part#, quantity, shipdate)

A relational table PARTSUPP has a composite primary key (supplier#, part#, shipdate)

Assume that:

- (i) a relational table PARTSUPP occupies 5000 data blocks,
- (ii) a blocking factor in a relational table PARTSUP is 100 rows per block,
- (iii) a relational table PARTSUPP contains information about 100 suppliers,
- (iv) a relational table PARTSUPP contains information about 500 parts,
- (v) a primary key is automatically indexed,
- (vi) an attribute part# is indexed,
- (vii) all indexes are implemented as B*-trees with a fanout equal to 20,
- (viii) a leaf level of an index on attribute part# consists of 50 data blocks,
- (ix) a leaf level of an index on primary key consists of 700 data blocks.

For each one of the following queries briefly describe how the database system processes each query and estimate the total number of read block operations needed to compute each query.

- (1)

```
SELECT quantity
FROM PARTSUPP
WHERE supplier# = 7 AND part# = 1 AND shipdate = '01-DEC-2019';
```
 - (2)

```
SELECT quantity
FROM PARTSUP
WHERE part# = 100 OR shipdate > '01-JAN-2020';
```
 - (3)

```
SELECT part#, COUNT(*)
FROM PARTSUPP
GROUP BY part#;
```
 - (4)

```
SELECT supplier#, part#, quantity
FROM PARTSUPP
ORDER BY supplier#, part#;
```
 - (5)

```
SELECT COUNT(*)
FROM PARTSUP
WHERE quantity > 1000 AND shipdate > '01-JAN-2020';
```
 - (6)

```
SELECT *
FROM PARTSUP
WHERE part# = 12345;
```
-

Question 4**(6 marks)**

Consider the following SELECT statements.

- (1) SELECT C_NATIONKEY, COUNT(*)
FROM CUSTOMER
GROUP BY C_NATIONKEY;
- (2) SELECT C_NATIONKEY, C_ACCTBAL
FROM CUSTOMER
ORDER BY C_NATIONKEY, C_ACCTBAL
- (3) SELECT COUNT(C_PHONE)
FROM CUSTOMER;
- (4) SELECT C_NATIONKEY, SUM(C_ACCTBAL)
FROM CUSTOMER
GROUP BY C_NATIONKEY;
- (5) SELECT *
FROM CUSTOMER
WHERE C_NATIONKEY = 12345 AND C_NAME = 'JAMES'
- (6) SELECT C_NAME
FROM CUSTOMER
WHERE C_ACCTBAL =100;

- (1) Find the smallest number of indexes that improve performance of all queries listed above.

(3 marks)

- (2) For each query briefly explain how the indexes will be used to process a query.

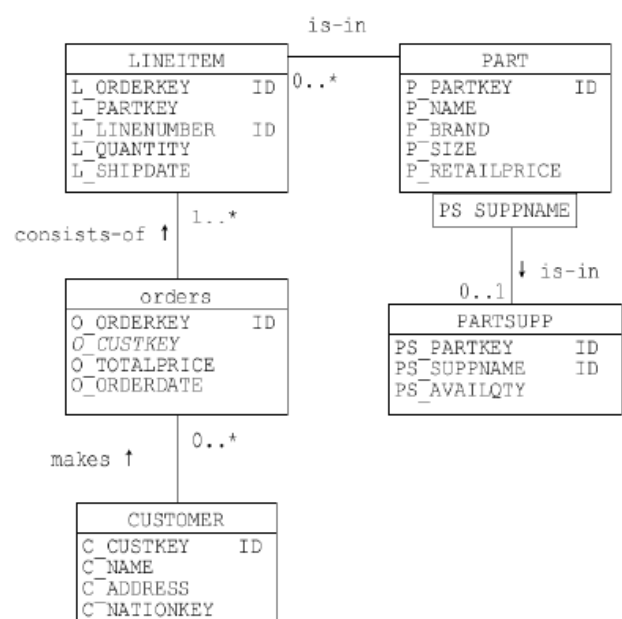
(3 marks)

Question 2

2. SELECT (SELECT COUNT(DISTINCT O_TOTALPRICE)
FROM ORDERS) PTOTAL,
(SELECT COUNT(DISTINCT O_ORDERDATE)
FROM ORDERS) DTOTAL
FROM DUAL;

Since both the attributes **o_totalprice** and **o_orderdate** are not found in any of the indexes, the query processor will perform a full-table scan on the table ORDERS. Retrieval can be improved if the following index is created:

create index ordersIdx1 on
ORDERS(o_orderdate, o_totalprice);



Question 5**(6 marks)**

Consider the following `SELECT` statements.

- (1)

```
SELECT C_CUSTKEY
FROM CUSTOMER
WHERE ( SELECT COUNT(*)
        FROM ORDERS
        WHERE ORDERS.O_CUSTKEY = CUSTOMER.C_CUSTKEY ) > 10;
```
- (2)

```
SELECT DISTINCT (SELECT COUNT(*)
                  FROM PART P
                  WHERE P.P_BRAND = PART.P_BRAND) TOTAL, P_BRAND
FROM PART;
```
- (3)

```
CREATE INDEX IDX ON PART(P_NAME);

SELECT *
FROM PART
WHERE (UPPER(P_NAME) = 'BOLT' AND P_RETAILPRICE > 2) ;

DROP INDEX IDX;
```
- (4)

```
SELECT O_ORDERKEY, O_CUSTKEY
FROM ORDERS
WHERE O_TOTALPRICE > 10
UNION
SELECT O_ORDERKEY, O_CUSTKEY
FROM ORDERS
WHERE O_TOTALPRICE < 5;
```

Find and write more efficient implementations of `SELECT` statements listed above.

Question 6**(8 marks)**

Consider a fragment of simple JDBC application listed below. It is a typical example of a pretty poor, from performance point of view, JDBC program. Rewrite a code written below to improve the performance of the application it is included in. There is no need to write the entire JDBC application.

Explain all details why your version of JDBC code is more efficient than the original one.

```
ResultSet rset1 = stmt1.executeQuery(
    "SELECT P_PARTKEY FROM PART ORDER BY P_NAME" );
long p_partkey = 0;
while ( rset1.next() )
{
    p_partkey = rset1.getInt(1);
    ResultSet rset2 = stmt2.executeQuery(
        "SELECT COUNT(*) FROM LINEITEM " +
        "WHERE L_PARTKEY = " + p_partkey );
    long total;
    while ( rset2.next() )
    {
        total = rset2.getInt(1);
        if (total >= 30 )
            System.out.println( p_partkey + "    " + total);
    }
}
```

End of Examination